

Dynamic Insertion Policy (DIP) :

→ LRU fails in thrashing workloads

→ Application types
 * Cache friendly
 * Scanning / Stream
 * Thrashing

→ Modules of replacement policy
 * Insertion * Eviction
 * Promotion

Thrashing: When the CPU tries to repeatedly access data whose size is larger than the cache size, the data gets pushed out of the cache even if the intent is to use it.

LRU Insertion Policy (LIP) \Rightarrow Bimodal Insertion Policy (BIP) \Rightarrow Dynamic Insertion Policy (DIP)

Cache replacement

→ Victim selection policy : Which line is evicted
 → Insertion policy : Where the incoming line is placed.

Example:

Traditional LRU:

Victim selection policy \rightarrow LRU position
 Insertion policy \rightarrow MRU position

LRU Insertion Policy (LIP):

- Victim selection policy is LRU
- Changes only in "Insertion Policy".
 - * Treat everything in LRU position
 - * Promote to MRU only if it is referenced while in LRU status.
 - * **But no ageing mechanism**

↳ So Bimodal Insertion Policy (BIP) is proposed.

Bimodal Insertion Policy (BIP):

→ controlled by Bimodal throttle parameter

- * Same as LIP but "with low probability" follows a MRU insertion policy.
- * It's like a middle ground between LRU & LIP.
- * Has an **ageing mechanism** while still maintaining **thrashing protection properties**.

↳ But still for an LRU friendly workload, DIP is problematic.

↳ **Dynamic Insertion Policy (DIP)**

is proposed.

Dynamic Insertion Policy:

- * Middle ground between traditional LRU and BIP.
- * Tries both LRU and BIP and checks which has the lower miss rate. And uses that.

↳ But this causes a lot of hardware overhead. So "Set Dueling" is proposed.

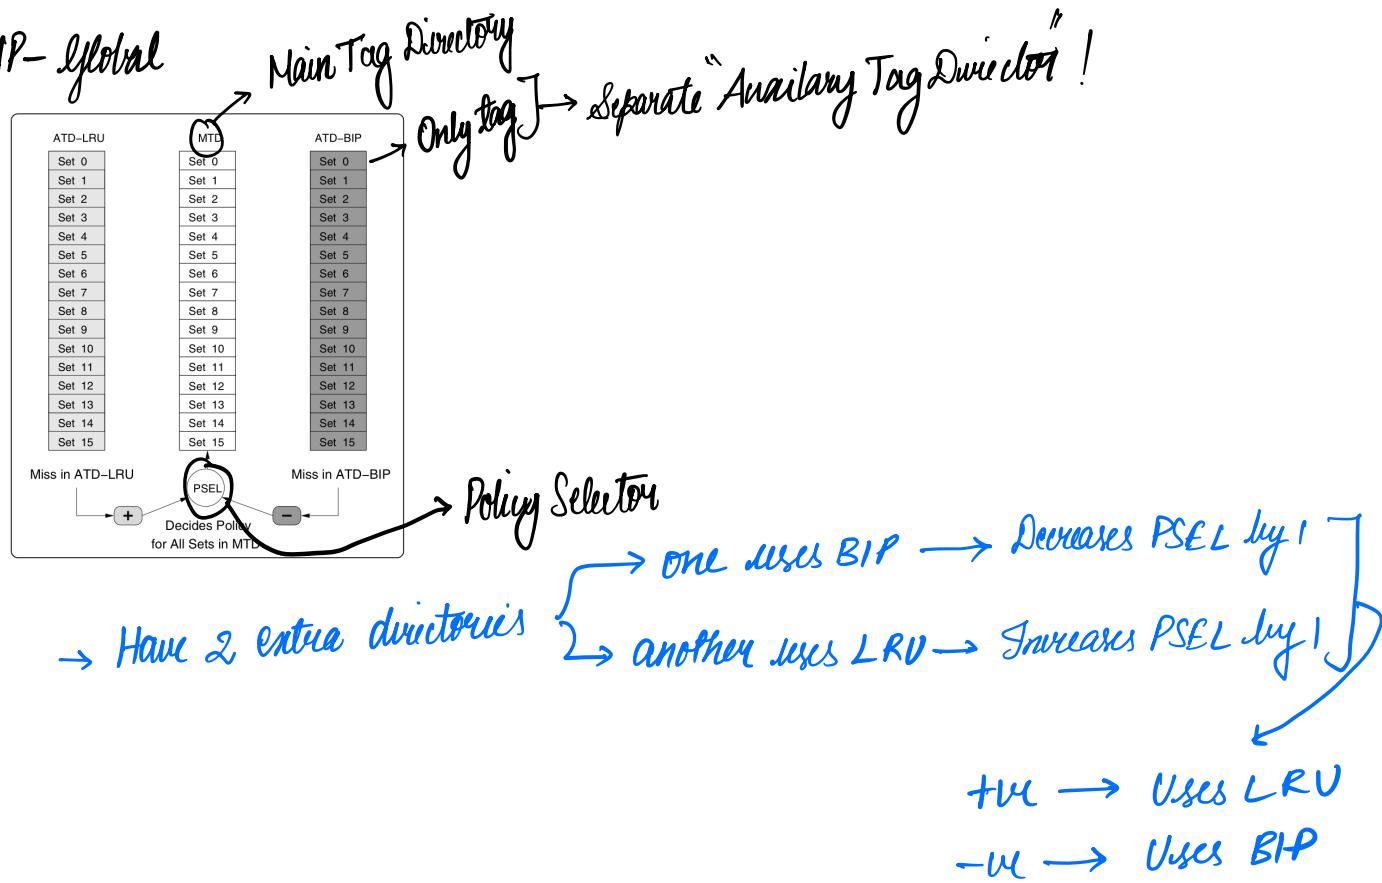
Zero-use lines: Lines not referenced between insertion and eviction

Doubt: "We do not enforce inclusion in our memory model." → Needs?

Doubt: "Demand references" → Send exact

Doubt: "Prefetcher pollution".

We only need tag to calculate the statistics

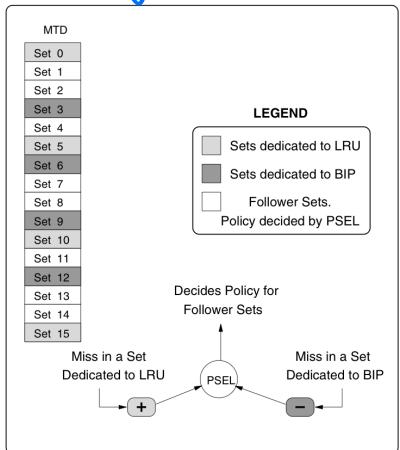


→ Should maintain 2 additional directories → Overhead!

→ To slightly reduce overhead **Dynamic Set Sampling** → cache behaviour can be approx.
↳ Still high. by a subset.

Proposes **Set Dueling**

Set dueling:



→ Dedicated sets

- * Few sets dedicated to LRU
- * Few for BIP

Instead have ATD
we are using the MTD
directly

→ Follower sets

→ But now the question of which sets should be chosen for BIP and LRU.



Dedicated Set selection Policy:

Statistically at design time

Dynamically at runtime

N sets
↓
Divide by K
↓
Each of the N/K region → called as constituency.

One set per constituency becomes a dedicated set for each of the policy.

Complement-Select Policy

The section introduces a complement-select policy, which simplifies the selection process:

- **Set index:** Every set in the cache is indexed using $\log_2(N)$ bits.
- Out of these bits, the most significant $\log_2(K)$ bits identify the constituency, while the remaining bits (the rest of the $\log_2(N/K)$ bits) identify the offset within the constituency.

The complement-select policy determines which sets are dedicated to which policies:

- LRU (Least Recently Used) policy gets sets where the constituency identifying bits are equal to the offset bits.
- BIP (Bimodal Insertion Policy) gets sets where the offset is the complement of the constituency identifying bits.

Example

For example, if you have a cache with 1024 sets and want to dedicate 32 sets to both LRU and BIP:

- The policy dedicates set 0 and every 33rd set to LRU.
- Set 31 and every 31st set would go to BIP.

Thus, LRU can be identified by comparing bits in the set index, and similarly, BIP can be identified by using the complement of those bits.

Summary

In this section, a complement-select policy is used to allocate sets between two competing policies, LRU and BIP, without requiring per-set leader markings. The dedicated sets are determined by comparing the constituency identifying bits with the offset within the constituency. This policy is explained in the context of how cache sets are divided and assigned under these competing policies.

→ Removes the need to separately remember which sets for which policy

K ↗ LRU → constituency identify bits = offset bits
K ↗ BIP → constituency identify bit complement offset bits

2. Constituency and Offset:

Let's say we are dividing the sets into **32 constituencies**. The number of bits needed to identify the constituency is $\log_2(32) = 5$ bits. This means the first 5 bits of the 10-bit set index represent the **constituency**.

The remaining 5 bits will represent the **offset** within that constituency. So, each constituency contains $N/K = 1024/32 = 32$ sets.

3. Equal and Complement Rules:

The **equal** and **complement** rules determine how sets are dedicated to different policies:

- **Equal:** The LRU (Least Recently Used) policy will get all sets where the **offset** bits are equal to the **constituency** bits.
- **Complement:** The BIP (Bimodal Insertion Policy) will get all sets where the **offset** is the complement of the **constituency** bits.

Example of the Equal Rule:

Let's take **constituency 3**. The binary representation of 3 is **00011**.

- The **equal rule** would dedicate any set where the **offset** is also **00011** (binary representation of 3) to the **LRU policy**.
- Therefore, set **0001100011** (which is set 99 in decimal) would be dedicated to LRU because its **constituency bits** (00011) are equal to its **offset bits** (00011).

Example of the Complement Rule:

Now let's look at the **complement rule** for **BIP**. The **complement** of a binary number is simply flipping all the bits (0s become 1s and vice versa).

- For **constituency 3** (which is **00011**), the complement of **00011** is **11100**.
- So, any set where the **offset** is **11100** would be dedicated to **BIP**.
- Therefore, set **0001111100** (which is set 124 in decimal) would be dedicated to BIP because its **constituency bits** (00011) are the complement of its **offset bits** (11100).

Applying It in a Real Example:

- Let's say we have a cache with **1024 sets** and we're dividing them into **32 constituencies**.
- For **constituency 3** (binary 00011):
 - The **equal rule** dedicates set **99** (0001100011) to **LRU**.
 - The **complement rule** dedicates set **124** (0001111100) to **BIP**.

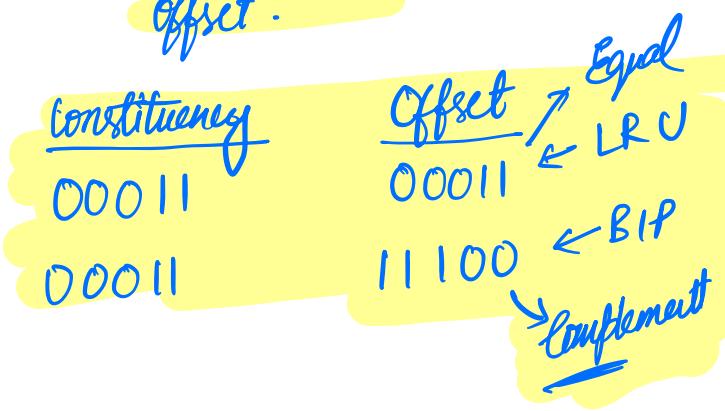
Thus, the **equal rule** means the **offset bits** match the **constituency bits**, and the **complement rule** means the **offset bits** are the bitwise complement of the **constituency bits**.

Summary:

- **Equal Rule:** Assigns a set to a policy if the **offset bits** are the same as the **constituency bits**.
 - Example: Constituency 3 (00011) → Set 99 (0001100011) is assigned to LRU.
- **Complement Rule:** Assigns a set to a policy if the **offset bits** are the **complement** of the **constituency bits**.
 - Example: Constituency 3 (00011) → Set 124 (0001111100) is assigned to BIP.

→ Example

→ If there are $\frac{N}{K} = 32$ constituency
then $\log_2 32 = 5$ bits are
constituency bits & remaining are
offset.



Mindmap:

